

Real Transfer Signal Brief

Synthetic conflict probes on production prompts

April 24, 2026. Medium validation run wr_14f78308dbac_dbc78513.

HEADLINE

Synthetic probes trained on clean policy-source conflict transfer to real DX Terminal prompts at L32 settings_end. trade_size fires most strongly on current-prefix concrete sized-action conflict: wrong buy/sell side, wrong allocation, or explicit size mismatch. shared_mean captures a broader policy-tension component, while both directions fire less strongly on rows whose evidence is temporal, bookkeeping-based, or interpretive.

Setup

Medium corpus: dx_terminal_signal_discovery_phase13_v1 in Neon. Scope: aggressive tier only, end-of-section sites, 500 complaints, 300 structure-matched controls, 118 Stage 1b strict anchors, and 33 buy-only anchors. Primary preregistered cell: layer 32, position settings_end. This is late enough for settings to be visible and close to the synthetic probe discovery band.

Definitions

Term	Meaning in this brief
probe / direction	fixed vector learned from synthetic conflict prompts; real prompts are scored by projection onto it
trade_size	probe family for concrete buy/sell/position-size disagreement
shared_mean	mean of the validated family directions; intended to capture common policy-source tension
L32 settings_end	model layer 32 activation at the end of the ACTIVE SETTINGS section
anchor positive	bridge rows designed to contain known conflict signal
structure-matched control	non-complaint production prompt with similar DX Terminal prompt shape
top / bottom	highest complaint projections / complaints closest to the control mean

The run did not train classifiers or thresholds. It only asked whether fixed synthetic directions produce meaningful scalar structure on real prompts.

Cohort Result

Direction	Anchor	Complaint	Control	Anchor-control	Complaint-control
trade_size	4.425	3.803	3.278	+1.147	+0.526
shared_mean	3.462	3.137	2.760	+0.703	+0.377

Interpretation: the transfer signal replicated at L32 settings_end for trade_size and shared_mean. Family specificity still matters: risk_preference was weaker at this cell and diversification_preference did not show clean ordering, so this is not every synthetic vector lighting up on every complaint prompt.

Why The First Proxy Was Too Coarse

The preregistered category proxy was USER_CONFIG_CONFLICT high and RULE_FABRICATION low. That is not the right target for trade_size. Those labels diagnose root cause; the probe target is visible conflict shape.

CORRECTION

Use action-shape labels for this audit: concrete current-prefix buy/sell/size conflict versus temporal, bookkeeping, fulfillment, or interpretation-layer conflict. Root-cause labels are metadata, not the primary semantic target.

Top/Bottom Readout

What The Probe Ranked High

Top/bottom slices were pulled from complaint rows only.

Direction	Top action/size	Top strategy	Bottom action/size	Bottom strategy
trade_size	20/25	5/25	15/25	10/25
shared_mean	20/25	5/25	9/25	16/25

Top trade_size complaint types: UNWANTED_BUY 10, UNWANTED_SELL 6, WRONG_SIZE 4. The high shared_mean slice is the same row set as high trade_size, but with lower projection magnitude. The clearest contrast is low shared_mean, where diffuse STRATEGY_IGNORED rows dominate.

Shape Of Low-Projection Conflicts

The bottom rows are not “agent did nothing.” Reviewed strategy-ignored bottom rows still contain buy/sell actions. Their conflict is different:

- no active strategy is visible, but the complaint references an old or expected strategy
- the current decision is already taking the requested action, so the complaint is retrospective
- the current action is an intermediate step in a multi-step strategy
- the disagreement is about entry price, timing, strategy fulfillment state, or rule interpretation rather than current visible trade size

Slice	Complaint shape	Visible prompt/action shape
high trade_size	wrong buy/sell or explicit size complaint	current-prefix action/size tension
high shared_mean	same high row set as trade_size	broader policy tension on concrete action conflicts
low trade_size	missed entry, old strategy, “why not execute?”	no active strategy, aligned action, or history-dependent conflict
low shared_mean	strategy ignored / rule narrative	temporal or interpretation-layer evidence dominates

Claim And Next Move

CLAIM

Fixed synthetic directions recover a real production signal at L32 settings_end. trade_size is selective for current-prefix concrete sized-action conflict. shared_mean tracks broader policy tension, while firing less on temporal, bookkeeping, and interpretation-layer complaints.

Residual risk: top/bottom labels are still an audit over projection extremes, not a new gold dataset. Recommended next move: hand-label the same rows with current_action_size_conflict, retrospective_history_conflict, strategy_fulfillment_conflict, interpretation_or_rule_conflict, and unclear_or_label_mismatch; then report enrichment against those labels before expanding cells or moving toward causal tests.