

Conflict Probe v0

Workflow report generated 2026-04-02T19:46:45.991026+00:00. This report summarizes the latest successful dataset, capture, and analysis chain for the selected workflow spec.

SPEC	PUBLICATION	ROWS	ANALYSIS RUN
conflict_probe_v0	workflow_dataset_conflict_probe_v0_v1	175	2ed5cbe4ed55

Spec	Publication	Rows
Conflict Probe v0	workflow_dataset_conflict_probe_v0_v1	175

Workflow

- Spec id: conflict_probe_v0
- Version: 1
- Dataset run: 1bdf3bed5b62
- Capture run: 9c0467b2ea48
- Analysis run: 2ed5cbe4ed55

Analysis Setup

- Mode: probe
- Target: workflow_label
- Data source: residual
- Pooling: last_token
- Labels parquet: data/workflows/conflict_probe_v0/labels/workflow_dataset_conflict_probe_v0_v1.parquet
- Output dir: research/prompt_confusion/outputs/conflict_probe_v0_analysis/2ed5cbe4ed55

Executive Read

MAIN READ

The best probe layer does not beat the majority baseline, so this run does not support a useful linear readout for the workflow label.

Main Quantitative Findings

Best Layer	Best Accuracy	Majority Baseline	Lift vs Baseline
34	0.576	0.627	-0.051

Why this matters:

- best-layer balanced accuracy is 0.524

- run covers 48 probed layers
- best-layer fold std is 0.051

Top Layers

Layer	Accuracy	Std	Balanced	Selectivity
34	0.576	0.051	0.524	0.035
6	0.525	0.073	0.500	0.037
31	0.523	0.040	0.472	0.016
39	0.517	0.046	0.486	-0.048
41	0.515	0.035	0.471	-0.016
5	0.509	0.059	0.507	0.024
4	0.504	0.039	0.493	0.005
38	0.504	0.058	0.470	0.032

Artifacts

- Results JSON: research/prompt_confusion/outputs/conflict_probe_v0_analysis/2ed5cbe4ed55/results.json
- Primary parquet: research/prompt_confusion/outputs/conflict_probe_v0_analysis/2ed5cbe4ed55/probe_workflow_label_residual.parquet
- Remote output: /data/analysis_results/workflows/conflict_probe_v0/2ed5cbe4ed55